

FPGA 编译项目报告

——技术报告

高琦，苏灿，黎睿正

1. 项目背景

目前大模型部署存在的难点

- 工具结合度低

现有的量化、打分和部署等工具模块间甚至模块内操作逻辑不一致，许多操作之间仍需手动调整参数。

- 新模型适配成本高

现有的工作流需要根据所使用的模型手动调整，新模型需要逐一环节调试，无法开箱即用。

- 操作学习成本高

现有的工具多为命令行工具，缺少直观的交互逻辑和图形化界面，用户操作上手难度较大。

2. 项目内容

2.1 技术特性

- 量化：支持 GPTQ 指定 bits、group_size、desc_act
- 打分
 - 支持 lm-evaluation-harness 测试项
arc_easy、arc_challenge、gsm8k_cot、gsm8k_platinum_cot、hellaswag、mmlu、gpqa、boolq、openbookqa
 - 支持 EvalPlus 测试项
humaneval、mbpp
- GPU 部署：支持 vLLM 指定上下文长度、显存占用限制、服务端口、API 密钥
- 权重处理：支持 Compiler-VCU128 指定 bits、group_size、desc_act

- FPGA 服务：支持 Fast API 指定上下文长度、生成温度、服务端口、API 密钥
- GPU、Port 调度器：支持 GPU 设备图分类、调整调度显存占用量、设备锁机制
- WebUI：支持指定服务地址、用户管理

2.2 WebUI

前端 WebUI 主要支持用户进行一键式操作，为大模型上板做好准备工作，提供用户登录、各项部署细节选择、取消进程、实时日志查看服务。

- 网页界面展示

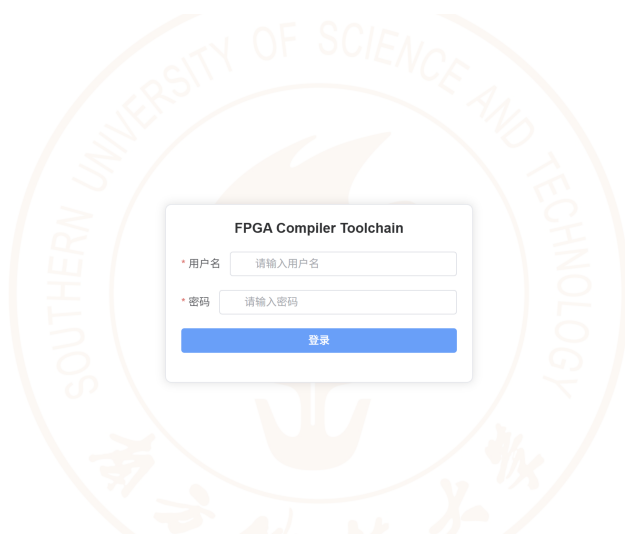


图 1 用户登录界面展示

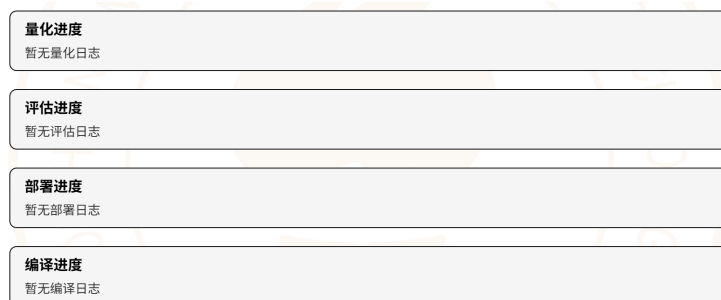


图 2 日志界面展示

 大模型 FPGA 部署工具

选择模型

请选择模型

选择量化精度

INT4 (仅支持)

选择评分框架

EvalPlus

选择评分任务

请选择评分任务

选择评分对象

原模型

开始部署

图 3 部署工具界面展示

```
//工具选项设置代码
data() {
  return {
    activeSections: ['model', 'quant', 'target', 'eval', 'task'],
    selectedModel: '',
    selectedQuantPrecision: 'int4',
    isDeploying: false,
    deployStatus: [],
    pollingInterval: null,
    progressPollingInterval: null,
    quantPid: null,
    selectedEvalMethod: 'evalPlus',
    selectedEvalTarget: 'none',
    quantLogs: [],
    evalLogs: [],
    dataLogs: [],
    models: [
      { value: 'qwen2', label: 'Qwen2-7B-Instruct', icon: modelQwen },
      { value: 'qwen2.5', label: 'Qwen2.5-7B-Instruct', icon: modelQwen },
      { value: 'qwen2-vl', label: 'Qwen2-VL-7B-Instruct', icon:
```

```

        modelQwen },
    { value: 'qwen2.5-vl', label: 'Qwen2.5-VL-7B-Instruct', icon:
        modelQwen },
    { value: 'deepseek', label: 'DeepSeek-R1-Distill-Qwen-7B', icon
        : modelDeepseek }
],
precisions: [
    { value: 'int2', label: 'INT2', precisionValue: 2 },
    { value: 'int4', label: 'INT4 (仅支持)', precisionValue: 4 },
    { value: 'int8', label: 'INT8', precisionValue: 8 }
],
evalMethods: [
    { label: 'EvalPlus', value: 'evalPlus' },
    { label: 'lmEvaluationHarness', value: 'lmEvalHarness' }
],
evalPlusTasks: [
    { value: 'humaneval', label: 'HumanEval' },
    { value: 'mbpp', label: 'MBPP' }
],
lmEvalHarnessTasks: [
    { value: 'arc_easy', label: 'ARC Easy' },
    { value: 'arc_challenge', label: 'ARC Challenge' },
    { value: 'gsm8k_cot', label: 'GSM8K CoT' },
    { value: 'gsm8k_platinum_cot', label: 'GSM8K Platinum CoT' },
    { value: 'hellaswag', label: 'HellaSwag' },
    { value: 'mmlu', label: 'MMLU' },
    { value: 'gpqa', label: 'GPQA' },
    { value: 'boolq', label: 'BoolQ' },
    { value: 'openbookqa', label: 'OpenBookQA' }
],
selectedEvalTasks: [],
evalTargets: [
    { label: '原模型', value: 'origin' },
    { label: '量化模型', value: 'quant' },
    { label: '两个都评分', value: 'both' },
    { label: '不评分', value: 'none' }
],
}
},

```

- 用户管理：支持自定义登录用户

```
//前端登录界面代码
```

```

<script>
import axios from 'axios'

export default {
  name: 'Login',
  data() {
    return {
      loading: false,
      loginForm: {
        username: '',
        password: ''
      },
      rules: {
        username: [
          { required: true, message: '请输入用户名', trigger: 'blur' }
        ],
        password: [
          { required: true, message: '请输入密码', trigger: 'blur' }
        ]
      }
    }
  },
  methods: {
    handleLogin() {
      this.$refs.loginForm.validate(valid => {
        if (valid) {
          this.loading = true
          this.verifyCredentials()
        }
      })
    },
    async verifyCredentials() {
      try {
        const authHeader = 'Basic ' + btoa(`${this.loginForm.username}:${this.loginForm.password}`)
        const response = await axios.get('/api/verify', {
          headers: { 'Authorization': authHeader }
        })

        if (response.data.success) {
          this.$emit('login-success', {
            username: this.loginForm.username,
            password: this.loginForm.password
          })
        }
      } catch (error) {
        console.log(error)
      }
    }
  }
}

```

```

    })
  } else {
    throw new Error(response.data.message || '认证失败')
  }
} catch (error) {
  let errorMsg = '登录失败: '
  if (error.response) {
    errorMsg += error.response.data?.message || error.response.
      statusText
  } else {
    errorMsg += error.message
  }
  this.$message.error(errorMsg)
  console.error('登录错误详情:', error)
} finally {
  this.loading = false
}
}
}
</script>

```

提供用户登录界面，使用体验更规范，同时可根据不同账户指定特定权限如管理员权限。支持客户添加指定账户，后续可拓展使用实时 token 验证方式提高安全性。

- 打分：支持同时开展多框架多测试项测试和原模型、量化模型对比测试

```

//轮询日志同步进度信息
async startEvaluationPolling(target) {
  let failCount = 0;
  const MAX_FAILS = 5;

  const interval = setInterval(async () => {
    try {
      const response = await axios.get(`${this.apiUrl}/eval_progress`,
        {
          headers: {
            'Authorization': 'Basic ' + btoa(`${this.authInfo.username}:${this.authInfo.password}`)
          }
        }
      );
    } catch (error) {
      failCount++;
      if (failCount === MAX_FAILS) {
        console.error('轮询失败');
        return;
      }
    }
  });

  if (response.data.success) {
    console.log('轮询成功');
  }
}

```

```

failCount = 0;
const progressLines = response.data.progress || [];
this.evalLogs.push(...progressLines);
this.$emit('eval-log', progressLines);

const hasError = this.evalLogs.some(line =>
  line.includes('[ERROR]') ||
  line.includes('失败') ||
  line.includes('异常') ||
  line.includes('Traceback')
);

const hasCompleted = this.evalLogs.some(line =>
  line.includes('完成') ||
  line.toLowerCase().includes('evaluation finished') ||
  line.toLowerCase().includes('scoring complete') ||
  line.toLowerCase().includes('done')
);

if (hasError) {
  clearInterval(interval);
  this.deployStatus.push(` ${target === 'origin' ? '原模型' : '
    量化模型'} 评分失败，请检查日志`);
  return;
}

if (!response.data.is_running && !hasCompleted) {
  clearInterval(interval);
  this.deployStatus.push(` ${target === 'origin' ? '原模型' : '
    量化模型'} 评分中断但未检测到“完成”关键词，可能失败`);
  return;
}

if (hasCompleted) {
  clearInterval(interval);
  this.deployStatus.push(` ${target === 'origin' ? '原模型' : '
    量化模型'} 评分完成`);
}
} catch (error) {
  const isAxiosError = error.isAxiosError;
  const errMsg = error?.message || '';
  const isIgnorable = errMsg.includes('ERR_EMPTY_RESPONSE') ||

```

```

        (isAxiosError && !error.response);

    if (isIgnorable) {
        failCount++;
        console.warn(` 评分轮询失败（可忽略）：${errMsg}, 当前失败次数：${failCount}`);

        if (failCount >= MAX_FAILS) {
            clearInterval(interval);
            this.deployStatus.push(` ${target === 'origin' ? '原模型' : '量化模型'} 连续多次无法获取评分进度，任务可能失败`);
            this.$message.error('评分进度查询连续失败，已中止');
        }
        return;
    }
    clearInterval(interval);
    this.deployStatus.push(` ${target === 'origin' ? '原模型' : '量化模型'} 评分进度获取失败：${errMsg}`);
    this.$message.error('评分进度查询失败');
}
}, 3000);
},

```

当流程进行到原模型打分或量化模型评分时调用 startEvaluation 启动评分程序

```

//评分部分开始
async startEvaluation(target) {
    const model = this.getCurrentModel();
    const method = this.selectedEvalMethod;

    this.evalLogs = [];
    this.deployStatus.push(`开始对 ${target === 'origin' ? '原模型' : '量化模型'} 进行评分（方法：${method}）...`);

    try {
        const response = await axios.post(`${this.apiUrl}`, {
            model_name: model.label,
            eval_method: method,
            eval_tasks: this.selectedEvalTasks,
            start_evaluation: true,
            is_quantized: target !== 'origin'
        }, {
            headers: {
                'Authorization': 'Basic ' + btoa(`${this.authInfo.username}:${this.authInfo.password}`)
            }
        });
    } catch (error) {
        console.error('评分失败：', error);
    }
}

```



```

        this.authInfo.password}`)
    }
  });

  if (response.data.success) {
    this.deployStatus.push(` ${target} === 'origin' ? '原模型' : '量化模型' } 评分任务已启动`);
    this.startEvaluationPolling(target);
  } else {
    throw new Error(response.data.message || '评分启动失败');
  }
} catch (error) {
  console.error(`评分启动失败 (${target})`, error);
  const errorMsg = error.response?.data?.message || error.message;
  this.deployStatus.push(` ${target} === 'origin' ? '原模型' : '量化模型' } 评分启动失败: ${errorMsg}`);
}
},

```

当用户取消部署时，前端向 api 发送取消信息，后端分别取消对应流程代码

```

//前端向api发送取消信息并监控是否成功取消进程
async cancelDeploy() {
  try {
    if (!this.isDeploying) return;

    if (this.progressPollingInterval) {
      clearInterval(this.progressPollingInterval);
    }

    this.deployStatus.push('正在取消部署流程...');

    try {
      const quantCancelResp = await axios.post(`${this.apiUrl}/cancel_quant`, {}, {
        headers: {
          'Authorization': 'Basic ' + btoa(`${this.authInfo.username}:${this.authInfo.password}`)
        }
      });
    }

    if (quantCancelResp.data.success) {
      this.deployStatus.push('已成功取消量化进程');
    } else {

```

```

        this.deployStatus.push(' 取消量化失败: ' + quantCancelResp.data.
            message);
    }
} catch (e) {
    this.deployStatus.push(' 取消量化时发生异常: ' + (e.message));
}

try {
    const cancelResp = await axios.post(`${this.apiUrl}/cancel_eval`,
        {}, {
            headers: {
                'Authorization': 'Basic ' + btoa(`${this.authInfo.username}:${this.authInfo.password}`)
            }
        });

    if (cancelResp.data.success) {
        this.deployStatus.push(` 已取消评分进程`);
    } else {
        this.deployStatus.push(` 无法取消评分: ${cancelResp.data.message} `);
    }
} catch (error) {
    this.deployStatus.push(` 取消评分失败: ${error.message}`);
}

try {
    const cancelResp = await axios.post(`${this.apiUrl}/
        cancel_deployment`, {}, {
            headers: {
                'Authorization': 'Basic ' + btoa(`${this.authInfo.username}:${this.authInfo.password}`)
            }
        });

    if (cancelResp.data.success) {
        this.deployStatus.push(` 已取消部署进程`);
        this.$message.success('部署进程已成功取消');
    } else {
        this.deployStatus.push(` 无法取消部署: ${cancelResp.data.message} `);
        this.$message.warning(cancelResp.data.message || '无法取消部署进程');
    }
}

```

```

    }
  } catch (error) {
    const errMsg = error?.response?.data?.message || error.message;
    this.deployStatus.push(` 取消部署失败: ${errMsg}`);
    this.$message.error('取消部署失败');
  }

  try {
    const cancelResp = await axios.post(`${this.apiUrl}/
      cancel_compile`, {}, {
        headers: {
          'Authorization': 'Basic ' + btoa(`${this.authInfo.username}:${this.authInfo.password}`)
        }
      });

    if (cancelResp.data.success) {
      this.deployStatus.push(` 已取消编译进程`);
    } else {
      this.deployStatus.push(` 无法取消编译: ${cancelResp.data.message}`);
    }
  } catch (error) {
    this.deployStatus.push(` 取消编译失败: ${error.message}`);
  }

  this.isDeploying = false;
  this.$message.warning('部署流程和评分流程已中断');

} catch (error) {
  console.error('取消部署失败:', error);
  const errorMsg = error.response?.data?.message || error.message;
  this.deployStatus.push(` 取消失败: ${errorMsg}`);
  this.$message.error(`取消失败: ${errorMsg}`);
} finally {
  this.isDeploying = false;
}
},

```

```

//取消评分进程
@app.route('/api/cancel_eval', methods=['POST'])
@auth.login_required
def cancel_evaluation():

```

```

global current_eval_process

try:
    if not is_eval_running():
        return jsonify({
            'success': False,
            'message': '没有正在运行的评估进程'
        }), 400

    with open(EVALUATION_LOG, 'a') as f:
        f.write("[INFO]正在取消评估进程...\n")

    current_eval_process.terminate()
    current_eval_process.join(timeout=2)

    with open(EVALUATION_LOG, 'a') as f:
        f.write("[INFO]评估进程已被用户取消\n")

    current_eval_process = None

    return jsonify({
        'success': True,
        'message': '评估进程已成功取消'
    })
except Exception as e:
    error_msg = f"取消评估失败:{str(e)}"
    log_error(error_msg, "backend")
    return jsonify({
        'success': False,
        'message': error_msg
    }), 500

```

- GPU 部署：支持同时部署原模型、量化模型以便对比

```

//根据用户需求进行对应模型部署
async startDeployment() {
    const model = this.getCurrentModel();

    this.deployLogs = [];
    this.deployStatus.push(`开始部署模型 ${model.label} ...`);

    try {
        const response = await axios.post(`${this.apiUrl}`, {

```

```

        model_name: model.label,
        start_deployment: true
      }, {
        headers: {
          'Authorization': 'Basic ' + btoa(`${this.authInfo.username}:${this.authInfo.password}`)
        }
      });

      if (response.data.success) {
        this.deployStatus.push(` 模型 ${model.label} 部署任务已启动`);
        this.startDeploymentPolling();
      } else {
        throw new Error(response.data.message || '部署启动失败');
      }
    } catch (error) {
      console.error('部署启动失败', error);
      const errorMsg = error.response?.data?.message || error.message;
      this.deployStatus.push(` 模型 ${model.label} 部署启动失败: ${errorMsg}`);
    }
  },

```

用户可在模型部署界面选择对原模型/量化模型进行部署，后续将补充对话框，实现部署后用户即时与模型互动直观体验功能

```

//模型部署界面
<template>
  <el-card class="deployment-tool-card">
    <div slot="header" class="header-with-icon">
      
      <span>部署模型管理</span>
    </div>

    <el-form label-position="top">
      <!-- 模型名称选择 -->
      <el-form-item label="模型名称">
        <el-select v-model="selectedModel" placeholder="请选择模型">
          <el-option
            v-for="model in models"
            :key="model.value"
            :label="model.label"
            :value="model.value">

```

```

        <span style="float: left">{{ model.label }}</span>
        
      </el-option>
    </el-select>
  </el-form-item>

  <!-- 部署类型选择 -->
  <el-form-item label="部署类型">
    <el-select v-model="deployType" placeholder="请选择部署类型">
      <el-option
        v-for="type in deployTypes"
        :key="type.value"
        :label="type.label"
        :value="type.value">
        <span style="float: left">{{ type.label }}</span>
      </el-option>
    </el-select>
  </el-form-item>

  <!-- 启动部署按钮 -->
  <div class="deploy-button-wrapper">
    <el-form-item>
      <el-button type="primary"
        :loading="isDeploying"
        :disabled="!selectedModel || !deployType"
        @click="startDeploy">
        启动部署
      </el-button>
    </el-form-item>
  </div>

  <!-- 部署日志展示区域 -->
  <el-card class="deploy-log-card" v-if="deployLogs.length > 0">
    <div class="log-title">部署日志</div>
    <div class="log-content">
      <div v-for="(log, index) in deployLogs" :key="index" class="
        log-line">{{ log }}</div>
    </div>
  </el-card>

</el-form>
</el-card>
</template>

```

- API 接口：捕捉用户需求并启动后端程序

```
WORKSPACE_ROOT = "/data/disk0/Workspace/Compiler-Toolchain/Compiler-Toolchain"
sys.path.insert(0, WORKSPACE_ROOT)
os.chdir(WORKSPACE_ROOT)
//以量化板块为例展示进程启动与日志捕获
def quantification_entrypoint(model_id, log_path, is_vl_model):
    try:
        gptq_log_dir = os.path.join(WORKSPACE_ROOT, "CT", "WebUI", "gptq_log")
        os.makedirs(gptq_log_dir, exist_ok=True)
        os.chdir(gptq_log_dir)

        with open(log_path, 'a') as f:
            with contextlib.redirect_stdout(f), contextlib.redirect_stderr(f):
                if is_vl_model:
                    from CT.Example.Quantization.qwenVLQuantization import simpleQuantization
                    print(f"[INFO] 启动VL模型量化: {model_id}")
                else:
                    from CT.Example.Quantization.quantization import simpleQuantization
                    print(f"[INFO] 启动普通模型量化: {model_id}")

                simpleQuantization(model_id)

                print(f"[INFO] 模型量化完成: {model_id}")

    except Exception as e:
        with open(log_path, 'a') as f:
            f.write(f"[ERROR] 模型量化异常: {e}\n")

def is_quant_running():
    global current_quant_process
    return current_quant_process is not None and current_quant_process.is_alive()

def run_quantification(model_name):
    global current_quant_process
```

```

try:
    # 清空进度日志
    with open(PROGRESS_LOG, 'w') as f:
        f.write("")

    is_vl_model = "VL" in model_name.upper()
    log_error(f"开始量化模型:{model_name}(类型:{'VL' if is_vl_model else '普通'}", "quant")

    # 创建并启动量化进程
    current_quant_process = multiprocessing.Process(
        target=quantification_entrypoint,
        args=(model_name, PROGRESS_LOG, is_vl_model,)
    )
    current_quant_process.start()

    return current_quant_process.pid

except Exception as e:
    error_msg = f"量化失败-模型:{model_name}错误:{traceback.format_exc()}"
    log_error(error_msg, "quant")
    current_quant_process = None
    raise

```

```

//接收前端信息执行对应操作
@app.route('/api', methods=['POST'])
@auth.login_required
def post_api():
    try:
        if data.get("start_quantization"):
            if "model_name" not in data:
                error_msg = "缺少模型名称参数"
                log_error(error_msg, "backend")
                return jsonify({'success': False, 'message': error_msg}),
                    400

    try:
        if is_quant_running():
            current_quant_process.terminate()
            time.sleep(1)

        pid = run_quantification(data["model_name"])

```



```

        log_error(f"已启动量化进程PID:_{pid}", "backend")
        return jsonify({
            'success': True,
            'message': '量化进程已启动',
            'pid': pid,
            'current_params': {
                'model_name': data["model_name"]
            }
        })
    except Exception as e:
        error_msg = f"量化进程启动失败:_{str(e)}"
        log_error(error_msg, "backend")
        return jsonify({'success': False, 'message': error_msg}),
            500

@app.route('/api/progress', methods=['GET'])
@auth.login_required
def get_progress():
    """获取量化进度"""
    try:
        if not os.path.exists(PROGRESS_LOG):
            return jsonify({
                'success': False,
                'message': '进度文件不存在',
                'is_running': False
            }), 404

        with open(PROGRESS_LOG, 'r') as f:
            lines = f.readlines()[-150:]

        # 过滤 ANSI 转义字符
        clean_lines = [remove_ansi_codes(line.strip()) for line in
            lines if line.strip()]

        return jsonify({
            'success': True,
            'progress': clean_lines,
            'is_running': is_quant_running()
        })
    except Exception as e:
        log_error(f"获取进度失败:_{str(e)}", "backend")
        return jsonify({

```

```

        'success': False,
        'message': '获取进度失败',
        'is_running': False
    }), 500

@app.route('/api/cancel_quant', methods=['POST'])
@auth.login_required
def cancel_quantization():
    global current_quant_process

    try:
        if not is_quant_running():
            return jsonify({
                'success': False,
                'message': '没有正在运行的量化进程'
            }), 400

        with open(PROGRESS_LOG, 'a') as f:
            f.write("[INFO]正在取消量化进程...\n")

        current_quant_process.terminate()
        current_quant_process.join(timeout=2)

        with open(PROGRESS_LOG, 'a') as f:
            f.write("[INFO]量化进程已被用户取消\n")

        current_quant_process = None

        return jsonify({
            'success': True,
            'message': '量化进程已成功取消'
        })
    except Exception as e:
        error_msg = f"取消量化失败:{str(e)}"
        log_error(error_msg, "backend")
        return jsonify({
            'success': False,
            'message': error_msg
        }), 500

```

- 客制化：支持定制页面图标、企业名称、语言前端代码通过组件进行组合，针对客户指定的页面需求可快速调整背景图片、图标、网页风格等一系列元素，实

现便捷定制。

2.3 FPGA OpenAI-Style API

- OpenAI-Style API 提供了快捷的与 FPGA 部署模型对话的工具, 在前端 openwebui 或是 agent 直接通过 OpenAI api chat/completion 端口均可拉起工具
- openwebui 界面展示



图 4 FPGA 模型对话界面展示

- socket 实现与上位机的通信, 在指定端口发送和接收消息

```
class FPGAConnection:
def __init__(self):
    self.conn = None
    self.shutdown_event = threading.Event()
    self.lock = threading.Lock()

def connect(self):

    while not self.shutdown_event.is_set():
        try:
            print(f"[H100]尝试连接上位机_{TCP_HOST}:{TCP_PORT}")
            sock = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
            sock.settimeout(10)
            sock.connect((TCP_HOST, TCP_PORT))
            self.conn = sock
            print(f"[H100]成功连接到上位机")
            return True
        except Exception as e:
            print(f"[H100]连接失败:{str(e)}")
            time.sleep(RECONNECT_DELAY)
```

```

        return False

def send_message(self, message):

    if self.conn is None:
        raise ConnectionError("未连接到上位机")

    try:
        with self.lock:

            full_message = message + END_MARKER.decode('utf-8')
            encoded_msg = full_message.encode('utf-8')

            msg_length = len(encoded_msg)
            self.conn.sendall(struct.pack('>I', msg_length))

            self.conn.sendall(encoded_msg)
            print(f"[H100] 发送消息到上位机 (长度: {msg_length})")

            time.sleep(2)

            length_data = self.conn.recv(4)
            if not length_data:
                raise ConnectionError("连接已关闭")
            msg_length = struct.unpack('>I', length_data)[0]

            chunks = []
            bytes_received = 0
            while bytes_received < msg_length:
                chunk = self.conn.recv(min(msg_length -
                    bytes_received, BUFFER_SIZE))
                if not chunk:
                    raise ConnectionError("连接中断")
                chunks.append(chunk)
                bytes_received += len(chunk)

            full_message = b''.join(chunks).decode('utf-8')

```

```

        if full_message.endswith(END_MARKER.decode('utf-8')):
            full_message = full_message[:-len(END_MARKER)]

        print(f"[H100] 收到上位机响应 (长度: {len(full_message)})")
        return full_message

    except Exception as e:
        print(f"[H100] 通信错误: {str(e)}")
        self.close()
        raise

def close(self):
    """关闭连接"""
    if self.conn:
        try:
            self.conn.close()
        except:
            pass
        self.conn = None
    print("[H100] 连接已关闭")

```

- FPGA 回传消息通过转化层处理为标准 OpenAI api 格式，可供 agent 直接读取

```

def process_requests():
    global fpga_connection

    while True:
        try:
            queue_item = request_queue.get()

            if len(queue_item) == 2:
                _, data = queue_item
                response_callback = None
            elif len(queue_item) == 3:
                _, data, response_callback = queue_item
            else:
                print(f"[H100] 错误: 无效的队列项格式: {queue_item}")
                continue

            with connection_lock:
                if fpga_connection is None or fpga_connection.conn is

```

```

        None:
            print("[H100]␣错误：未连接上位机")
            continue

    try:
        with connection_lock:
            current_conn = fpga_connection

            send_data = {
                "model": data.get("model"),
                "messages": data.get("messages"),
                "temperature": data.get("temperature"),
                "max_tokens": data.get("max_tokens")
            }

            response = current_conn.send_message(json.dumps(send_data))

            openai_response = {
                "id": f"chatcmpl-{int(time.time())}",
                "object": "chat.completion",
                "created": int(time.time()),
                "model": openai_model,
                "choices": [{
                    "message": {
                        "role": "assistant",
                        "content": response
                    },
                    "finish_reason": "stop",
                    "index": 0
                }]
            }

            if response_callback:
                response_callback(openai_response)

    except Exception as e:
        print(f"[H100]␣处理请求失败：␣{str(e)}")

    with connection_lock:

```

```

        if fpga_connection:
            fpga_connection.close()
            fpga_connection = None

    except Exception as e:
        print(f"[H100]请求处理错误:{str(e)}")
        time.sleep(1)

```

- 与 agent 或网页进行收发消息交互

```

def openai_endpoint():

    try:

        with connection_lock:
            if fpga_connection is None or fpga_connection.conn is None:
                return jsonify({"error": "未连接上位机"}), 503

        request_data = request.json
        print(f"[H100]收到Agent请求:{json.dumps(request_data,indent
            =2)}")

        response_event = threading.Event()
        response_data = [None]

        def response_callback(resp):
            response_data[0] = resp
            response_event.set()

        processing_data = {
            "messages": request_data['messages'],
            "model": request_data.get('model', 'glm'),
            "max_tokens": request_data.get('max_tokens', 1024),
            "temperature": request_data.get('temperature', 0.7),
            "response_callback": response_callback
        }

        request_queue.put((1, request_data, response_callback))
        print("发送给agent回复")
        #request_queue.put(("fpga", processing_data))
        response_event.wait(timeout=60)

```

```

if response_data[0] is None:
    return jsonify({"error": "处理超时"}), 504

return jsonify(response_data[0])

except Exception as e:
    print(f"[H100] API错误: {str(e)}")
    return jsonify({"error": "服务器错误"}), 500

```

- 上位机监听服务器发送的消息，编译成 FPGA 可读的形式传给 FPGA

```

def other_server_communication():
    global other_server_socket
    server_socket = socket.socket(socket.AF_INET, socket.SOCK_STREAM)
    server_socket.setsockopt(socket.SOL_SOCKET, socket.SO_REUSEADDR, 1)
    server_socket.bind(("0.0.0.0", OTHER_SERVER_PORT)) # 监听所有接口
    server_socket.listen(1)
    print(f"上位机已在端口 {OTHER_SERVER_PORT} 监听H100连接...")
    while True:
        try:
            client_sock, addr = server_socket.accept()
            print(f"接受来自H100的连接: {addr}")
            other_server_socket = client_sock
            while True:

                length_data = other_server_socket.recv(4)
                if not length_data:
                    raise ConnectionError("连接已关闭")

                # 解析消息长度
                msg_length = struct.unpack('>I', length_data)[0]

                chunks = []
                bytes_received = 0
                while bytes_received < msg_length:
                    chunk = other_server_socket.recv(min(msg_length - bytes_received, 4096))
                    if not chunk:
                        raise ConnectionError("连接中断")
                    chunks.append(chunk)

```



```

        bytes_received += len(chunk)

    full_message = b''.join(chunks).decode('utf-8')

    if full_message.endswith(END_MARKER.decode('utf-8')):
        full_message = full_message[:-len(END_MARKER)]

    message_queue.put(full_message)
    print(f"收到完整服务器消息 (长度: {msg_length}): {
        full_message[:50]}...")

except (ConnectionError, socket.error) as e:
    print(f"服务器连接错误: {str(e)}")
    if other_server_socket:
        other_server_socket.close()
        other_server_socket = None
    time.sleep(5) # 等待后重连

except Exception as e:
    print(f"处理服务器消息出错: {str(e)}")
    if other_server_socket:
        other_server_socket.close()
        other_server_socket = None
    time.sleep(1)

def main_kvcache_show(tokenizer, client, memory=0):
    global other_server_socket

    server_thread = threading.Thread(target=
        other_server_communication, daemon=True)
    server_thread.start()

    round = 1
    while True:

        query, is_server_message = collect_input()

        if query is None:
            time.sleep(0.1)
            continue

        try:

```

```

data = json.loads(query)

user_messages = [msg for msg in data.get("messages", [])
                  if msg.get("role") == "user"]

if user_messages:

    actual_query = user_messages[-1]["content"]
else:

    system_messages = [msg for msg in data.get("messages", [])
                       if msg.get("role") == "system"]
    actual_query = system_messages[0]["content"] if system_messages else "没有提供查询内容"

temperature = data.get("temperature", [])
max_tokens = data.get("max_tokens", [])

print(f"解析后的查询内容: {actual_query}")
print(f"参数: temperature={temperature}, max_tokens={max_tokens}")

except json.JSONDecodeError:
    print("无法解析JSON, 使用原始输入")
    actual_query = query
except Exception as e:
    print(f"解析输入时出错: {str(e)}")
    actual_query = query

prompt = "[Round {}]\n\n问: {}\n\n答: ".format(round,
        actual_query)
print(f"构造的prompt: {prompt[:100]}...")
inputs = tokenizer([prompt], return_tensors="pt")

print("FPGA: ", end="")
run_model_kvcache_show(client, inputs["input_ids"], inputs["input_ids"].shape[1], 0, memory)

full_response = ""
while True:
    state, ids_len, ids = get_next_ids(client)

```

```

if state == 0:
    print("")
    main_time_show(client)

    if other_server_socket:
        send_to_other_server(full_response)
    break

    generated_text = tokenizer.decode(ids)
    print(generated_text, end="", flush=True)
    full_response += generated_text

round += memory

```

3. 开发计划表

- 未来会继续向现有的工具链中添加新的功能
- 除了在各个环节加入新的功能外，还将额外开发 Cli 命令行工具以便在服务器终端操作调用
- 苏灿同学会同时参与量化部分的 llmc 开发

子项目	开发者	已完成内容	6月29日	未来开发
量化	高琦	GPTQ	AWQ	VIT量化工具适配
打分		Lm-Evaluation-Harness EvalPlus	OpenCompass	llmc
GPU部署		vLLM	暂无	
权重处理	苏灿	Compiler-VCU128	多模态模型权重处理和上板	
API服务		Fast API	暂无	
调度器	高琦	GPU, Port	已完成	
WebUI	黎睿正	基础功能	适配后端GPU部署 测试功能	完成WebUI开发内容
Cli	高琦	暂不开发	适配完毕已有功能	

图 5 开发计划表